# Bios 6301: Assignment 1

## Max Rohde

### Create a Data Set

A data set in R is called a data.frame. This particular data set is made of three categorical variables, or factors: `gender`, `smoker`, and `exercise`. In addition `exercise` is an ordered factor. `age` and `los` (length of stay) are continuous variables.

```r
gender <- c('M','M','F','M','F','F','M','F','M')
age <- c(34, 64, 38, 63, 40, 73, 27, 51, 47)
smoker <- c('no','yes','no','no','yes','no','no','no','yes')
exercise <- factor(c('moderate','frequent','some','some','moderate','none','none','moderate','moderate')
                   levels=c('none','some','moderate','frequent'), ordered=TRUE
)
los <- c(4,8,1,10,6,3,9,4,8)
x <- data.frame(gender, age, smoker, exercise, los)
x
```

```
##   gender age smoker exercise los
## 1      M  34     no moderate   4
## 2      M  64    yes frequent   8
## 3      F  38     no     some   1
## 4      M  63     no     some  10
## 5      F  40    yes moderate   6
## 6      F  73     no     none   3
## 7      M  27     no     none   9
## 8      F  51     no moderate   4
## 9      M  47    yes moderate   8
```

### Create a Model

We can create a model using our data set. In this case I'd like to estimate the association between `los` and all remaining variables. This means `los` is our dependent variable. The other columns will be terms in our model.

The `lm` function will take two arguments, a formula and a data set. The formula is split into two parts, where the vector to the left of `~` is the dependent variable, and items on the right are terms.

```r
lm(los ~ gender + age + smoker + exercise, dat=x)
```

```
##
## Call:
## lm(formula = los ~ gender + age + smoker + exercise, data = x)
##
## Coefficients:
## (Intercept)      genderM         age    smokeryes   exercise.L   exercise.Q
##    0.588144     4.508675    0.033377     2.966623    -2.749852    -0.710942
##  exercise.C
##    0.002393
```

1. Looking at the output, which coefficient seems to have the highest effect on `los`? (2 points)

The gender coefficent seems to have the highest effect on `los` because the value is largest.

This can be tough because it also depends on the scale of the variable. If all the variables are standardized, then this is not the case.

Given that we only have nine observations, it's not really a good idea to include all of our variables in the model. In this case we'd be "over-fitting" our data. Let's only include one term, `gender`.

**Warning**   When choosing terms for a model, use prior research, don't just select the variable with the highest coefficient.

2. Create a model using `los` and `gender` and assign it to the variable `mod`. Run the `summary` function with `mod` as its argument. (5 points)

```
mod <- lm(los ~ gender, dat=x)

summary(mod)
```

```
##
## Call:
## lm(formula = los ~ gender, data = x)
##
## Residuals:
##    Min      1Q Median     3Q    Max
##   -3.8    -0.5    0.2    1.2    2.5
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)    3.500      1.099   3.186   0.0154 *
## genderM        4.300      1.474   2.917   0.0224 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.197 on 7 degrees of freedom
## Multiple R-squared:  0.5487, Adjusted R-squared:  0.4842
## F-statistic:  8.51 on 1 and 7 DF,  p-value: 0.02243
```

The summary of our model reports the parameter estimates along with standard errors, test statistics, and p-values. This table of estimates can be extracted with the `coef` function.

**Estimates**

3. What is the estimate for the intercept? What is the estimate for gender? Use the `coef` function. (3 points)

```
summary(mod) %>% coef()
```

```
##             Estimate Std. Error  t value    Pr(>|t|)
## (Intercept)      3.5   1.098701 3.185581 0.01537082
## genderM          4.3   1.474061 2.917110 0.02243214
```

The estimate for the intercept is 3.5
The estimate for gender is 4.3

4. The second column of `coef` are standard errors. These can be calculated by taking the `sqrt` of the `diag` of the `vcov` of the `summary` of `mod`. Calculate the standard errors. (3 points)

```r
summary(mod) %>% coef()
```

```
##             Estimate Std. Error  t value    Pr(>|t|)
## (Intercept)      3.5   1.098701 3.185581 0.01537082
## genderM          4.3   1.474061 2.917110 0.02243214
```

```r
summary(mod) %>% vcov %>% diag %>% sqrt
```

```
## (Intercept)     genderM
##    1.098701    1.474061
# We see that the standard error computed with this method
# agrees with those given by the `coef()` function
```

The third column of `coef` are test statistics. These can be calculated by dividing the first column by the second column.

```r
mod <- lm(los ~ gender, dat=x)
mod.c <- coef(summary(mod))
mod.c[,1]/mod.c[,2]
```

```
## (Intercept)     genderM
##    3.185581    2.917110
```

The fourth column of `coef` are p values. This captures the probability of observing a more extreme test statistic. These can be calculated with the `pt` function, but you will need the degrees-of-freedom. For this model, there are 7 degrees-of-freedom.

5. Use the `pt` function to calculate the p value for gender. The first argument should be the test statistic for gender. The second argument is the degrees-of-freedom. Also, set the `lower.tail` argument to FALSE. Finally multiple this result by two. (4 points)

```r
# Store the t-statistics for gender
t_stat_gender <- mod.c[,3][2]
pt(q=t_stat_gender, df=7, lower.tail = FALSE)*2
```

```
##    genderM
## 0.02243214
# The calculated p-value of 0.02243214 agrees with the value in the table
```

**Predicted Values**

The estimates can be used to create predicted values.

```r
3.5+(x$gender=='M')*4.3
```

```
## [1] 7.8 7.8 3.5 7.8 3.5 3.5 7.8 3.5 7.8
```

6. It is even easier to see the predicted values by passing the model `mod` to the `predict` or `fitted` functions. Try it out. (2 points)

```r
predict(mod)
```

```
##   1   2   3   4   5   6   7   8   9
## 7.8 7.8 3.5 7.8 3.5 3.5 7.8 3.5 7.8
```

```r
fitted(mod)
```

```
##   1   2   3   4   5   6   7   8   9
## 7.8 7.8 3.5 7.8 3.5 3.5 7.8 3.5 7.8
```

3

7. `predict` can also use a new data set. Pass `newdat` as the second argument to `predict`. (3 points)

```
newdat <- data.frame(gender=c('F','M','F'))

predict(mod, newdata = newdat)
```

```
##   1   2   3
## 3.5 7.8 3.5
```

**Residuals**

The difference between predicted values and observed values are residuals.

8. Use one of the methods to generate predicted values. Subtract the predicted value from the `x$los` column. (5 points)

```
x$los - predict(mod)
```

```
##    1    2    3    4    5    6    7    8    9
## -3.8  0.2 -2.5  2.2  2.5 -0.5  1.2  0.5  0.2
```

9. Try passing `mod` to the `residuals` function. (2 points)

```
residuals(mod)
```

```
##    1    2    3    4    5    6    7    8    9
## -3.8  0.2 -2.5  2.2  2.5 -0.5  1.2  0.5  0.2
```

10. Square the residuals, and then sum these values. Compare this to the result of passing `mod` to the `deviance` function. (6 points)

```
residuals(mod)^2 %>% sum()
```

```
## [1] 33.8
```

```
deviance(mod)
```

```
## [1] 33.8
```

Remember that our model object has two items in the formula, `los` and `gender`. The residual degrees-of-freedom is the number of observations minus the number of items to account for in the model formula.

This can be seen by passing `mod` to the function `df.residual`.

```
df.residual(mod)
```

```
## [1] 7
```

11. Calculate standard error by dividing the deviance by the degrees-of-freedom, and then taking the square root. Verify that this matches the output labeled "Residual standard error" from `summary(mod)`. (5 points)

```
(deviance(mod)/df.residual(mod)) %>% sqrt()
```

```
## [1] 2.197401
```

```
summary(mod)
```

```
##
## Call:
## lm(formula = los ~ gender, data = x)
##
## Residuals:
##     Min      1Q  Median      3Q     Max
```

4

```
##   -3.8   -0.5    0.2    1.2    2.5
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)    3.500      1.099   3.186   0.0154 *
## genderM        4.300      1.474   2.917   0.0224 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.197 on 7 degrees of freedom
## Multiple R-squared:  0.5487, Adjusted R-squared:  0.4842
## F-statistic:  8.51 on 1 and 7 DF,  p-value: 0.02243
# The Residual standard error is 2.197 using both methods
```

Note it will also match this output:

```
predict(mod, se.fit=TRUE)$residual.scale
```

```
## [1] 2.197401
```

**T-test**

Let's compare the results of our model to a two-sample t-test. We will compare `los` by men and women.

12. Create a subset of `x` by taking all records where gender is 'M' and assigning it to the variable `men`. Do the same for the variable `women`. (4 points)

```
dplyr::filter(x, gender=="M") -> men
dplyr::filter(x, gender=="F") -> women
```

13. By default a two-sampled t-test assumes that the two groups have unequal variances. You can calculate variance with the `var` function. Calculate variance for `los` for the `men` and `women` data sets. (3 points)

```
var_men <- var(men$los)
var_women <- var(women$los)

var_men
```

```
## [1] 5.2
```

```
var_women
```

```
## [1] 4.333333
```

14. Call the `t.test` function, where the first argument is `los` for women and the second argument is `los` for men. Call it a second time by adding the argument `var.equal` and setting it to TRUE. Does either produce output that matches the p value for gender from the model summary? (3 points)

```
t.test(women$los, men$los)
```

```
##
##  Welch Two Sample t-test
##
## data:  women$los and men$los
## t = -2.9509, df = 6.8146, p-value = 0.02205
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  -7.7647486 -0.8352514
## sample estimates:
```

```
## mean of x mean of y
##        3.5       7.8
```

```r
t.test(women$los, men$los, var.equal = TRUE)
```

```
##
##  Two Sample t-test
##
## data:  women$los and men$los
## t = -2.9171, df = 7, p-value = 0.02243
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  -7.7856014 -0.8143986
## sample estimates:
## mean of x mean of y
##        3.5       7.8
```

```r
# The t.test with var.equal=TRUE produces output that matches the p-value
# from the models summary
```

An alternative way to call `t.test` is to use a formula.

```r
t.test(los ~ gender, dat=x, var.equal=TRUE)
```

```
##
##  Two Sample t-test
##
## data:  los by gender
## t = -2.9171, df = 7, p-value = 0.02243
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  -7.7856014 -0.8143986
## sample estimates:
## mean in group F mean in group M
##             3.5             7.8
```

```r
# compare p-values
t.test(los ~ gender, dat=x, var.equal=TRUE)$p.value
```

```
## [1] 0.02243214
```

```r
coef(summary(lm(los ~ gender, dat=x)))[2,4]
```

```
## [1] 0.02243214
```