

# Bios 6301: Assignment 3

Max Rohde

## Question 1

### 15 points

Write a simulation to calculate the power for the following study design. The study has two variables, treatment group and outcome. There are two treatment groups (0, 1) and they should be assigned randomly with equal probability. The outcome should be a random normal variable with a mean of 60 and standard deviation of 20. If a patient is in the treatment group, add 5 to the outcome. 5 is the true treatment effect. Create a linear model for the outcome by the treatment group, and extract the p-value (hint: see assignment1). Test if the p-value is less than or equal to the alpha level, which should be set to 0.05.

Repeat this procedure 1000 times. The power is calculated by finding the percentage of times the p-value is less than or equal to the alpha level. Use the `set.seed` command so that the professor can reproduce your results.

1. Find the power when the sample size is 100 patients. (10 points)
2. Find the power when the sample size is 1000 patients. (5 points)

```
set.seed(27182)

simulate <- function(n){
  # Create the group assignment by using a random permutation
  groups <- sample(rep(c("Treatment", "Control"), n/2))

  # Create a data frame to organize our results
  df <- tibble(subject_id = 1:n,
               group = groups,
               outcome = rnorm(n, mean=60, sd=20))

  # Add 5 to the outcome for all subjects in the treatment group
  df <- mutate(df,
               outcome = ifelse(group=="Treatment", outcome+5, outcome))

  # Create the linear model
  model <- lm(outcome ~ group, data = df)

  # Get the p-value from the model
  p <- summary(model)$coefficients[[2,4]]

  # return TRUE if statistically significant, FALSE otherwise
  return(p<=0.05)
}

# Simulate 1000 times for studies of 100 subjects
power_n100 <- map_lgl(1:1000, ~simulate(n=100)) %>% mean()
```

```
# Simulate 1000 times for studies of 1000 subjects
power_n1000 <- map_lgl(1:1000, ~simulate(n=1000)) %>% mean()
```

The estimated power with 100 subjects is 0.238

The estimated power with 1000 subjects is 0.974

## Question 2

### 14 points

Obtain a copy of the football-values lecture. Save the 2020/proj\_wr20.csv file in your working directory. Read in the data set and remove the first two columns.

1. Show the correlation matrix of this data set. (4 points)

```
# Read in the data
df <- read_csv('https://raw.githubusercontent.com/couthcommander/football-values/master/2020/proj_wr20.csv')

# Remove first two columns
select(df, -(1:2)) -> df

# Show correlation matrix
cor(df)
```

```
##           rec_att  rec_yds  rec_tds  rush_att  rush_yds  rush_tds  fumbles
## rec_att  1.0000000 0.9915980 0.9762985 0.4243375 0.4208867 0.2955319 0.8599863
## rec_yds  0.9915980 1.0000000 0.9872802 0.4083053 0.4043565 0.2888207 0.8576713
## rec_tds  0.9762985 0.9872802 1.0000000 0.3949094 0.3902740 0.2773522 0.8433106
## rush_att 0.4243375 0.4083053 0.3949094 1.0000000 0.9813709 0.8502802 0.4507697
## rush_yds 0.4208867 0.4043565 0.3902740 0.9813709 1.0000000 0.8631215 0.4441662
## rush_tds 0.2955319 0.2888207 0.2773522 0.8502802 0.8631215 1.0000000 0.2763207
## fumbles  0.8599863 0.8576713 0.8433106 0.4507697 0.4441662 0.2763207 1.0000000
## fpts      0.9898311 0.9981916 0.9919078 0.4439299 0.4408144 0.3261509 0.8565599
##           fpts
## rec_att  0.9898311
## rec_yds  0.9981916
## rec_tds  0.9919078
## rush_att 0.4439299
## rush_yds 0.4408144
## rush_tds 0.3261509
## fumbles  0.8565599
## fpts      1.0000000
```

1. Generate a data set with 30 rows that has a similar correlation structure. Repeat the procedure 1,000 times and return the mean correlation matrix. (10 points)

```
# Store the column means and covariance matrix
means <- colMeans(df)
covariance <- var(df)

# This generates one data set, with 30 rows, that has similar correlation structure
# to the original dataset. We do this simulating a multivariate normal random variable
# with the same mean vector and covariance matrix
MASS::mvrnorm(30, mu=means, Sigma=covariance)
```

```
##           rec_att  rec_yds  rec_tds  rush_att  rush_yds  rush_tds
## [1,] 37.3852511 483.1540947 2.79949411 0.38775991 4.649269 0.068672804
```

```

## [2,] 7.1805103 151.8290694 1.34050199 1.21875884 6.697488 0.035633623
## [3,] -10.9329090 -215.5304309 -1.95013736 -2.90206291 -22.100596 -0.116676025
## [4,] 44.0808065 560.4871552 3.44138589 2.16638867 18.129144 0.030584587
## [5,] -41.6620466 -509.3530951 -3.16813259 -1.65906216 -15.049493 -0.017036452
## [6,] 33.9518038 465.1485947 3.46724161 5.00835995 29.334496 0.268234554
## [7,] -22.3212191 -289.8170074 -1.61920725 0.82536496 1.116605 -0.011520775
## [8,] 14.0773190 254.6279087 1.91698117 -0.83743397 -7.918677 -0.008679039
## [9,] 76.1828166 1008.4319956 6.41573097 2.83469984 22.302712 0.180778038
## [10,] -3.8734388 0.7773838 0.69316315 2.10813066 8.704510 0.151293976
## [11,] 47.1455815 624.2357370 3.83959253 3.52399424 21.451713 0.015001484
## [12,] 55.2322172 766.3196871 4.56688113 0.55616179 4.213652 0.047976621
## [13,] -17.0083748 -197.5274730 -1.37381356 1.52222918 14.572487 0.207736951
## [14,] 24.0664613 372.8099256 2.51078716 0.63594011 8.234054 0.010519562
## [15,] 2.9780645 4.5458340 0.06792477 0.31721760 2.982391 -0.003581532
## [16,] 55.7624594 730.8758636 4.52038187 0.44792281 1.374372 0.041887476
## [17,] -0.2654242 62.9972732 0.21564744 -2.39618154 -22.417075 -0.092271584
## [18,] 13.8922647 89.9633039 0.47886711 3.97598725 24.981426 0.116495461
## [19,] 26.7110558 422.0751013 2.66325940 1.70303724 11.265567 0.133594461
## [20,] 61.7661704 886.3340374 5.45732863 -0.48847547 -7.354597 -0.141422254
## [21,] 35.6399621 491.4669472 3.01658328 3.05197958 13.526096 -0.053687420
## [22,] 41.4112354 515.0334892 3.22471088 2.15833949 12.799373 0.061377369
## [23,] 14.7973357 195.8724085 1.30710526 -0.06547853 1.017905 -0.011853824
## [24,] -31.5716023 -361.6760655 -1.80811409 -2.10352114 -15.899716 -0.101673599
## [25,] 14.3837974 258.6776484 1.40138783 -3.04578172 -20.104777 -0.174815291
## [26,] 55.3891036 735.5899278 4.71920099 5.62216948 37.210784 0.128573993
## [27,] 16.8169043 204.3673431 1.17099905 2.91775236 23.108334 0.182047782
## [28,] 45.6635614 643.1927315 3.94199357 -3.03079217 -19.672453 -0.108952336
## [29,] 69.7585499 900.4064968 5.67739869 6.53154715 45.504903 0.287702830
## [30,] 30.2890135 337.0130405 2.20908319 -3.85569600 -30.801941 -0.210344215
## fumbles fpts
## [1,] 0.21165261 65.3561832
## [2,] 0.15116317 23.5907891
## [3,] -0.02936617 -36.0252998
## [4,] 0.41230902 77.8744809
## [5,] -0.40339247 -70.6279996
## [6,] 0.18152630 71.1903593
## [7,] -0.31800133 -38.1044252
## [8,] 0.21139789 35.6675025
## [9,] 0.61495097 141.6651113
## [10,] -0.12972804 6.1732997
## [11,] 0.59824412 86.2418889
## [12,] 0.48382797 104.1462353
## [13,] -0.05597154 -25.2138341
## [14,] 0.03440669 53.1945553
## [15,] 0.06768509 0.9068102
## [16,] 0.30465676 100.1299288
## [17,] -0.11645884 5.1685344
## [18,] 0.46694675 14.5078958
## [19,] 0.27142100 59.7369602
## [20,] 0.53640734 118.8303987
## [21,] 0.31837778 67.5277492
## [22,] 0.05875391 72.4458245
## [23,] -0.06379426 27.3093059
## [24,] -0.39895065 -48.6674510

```

```
## [25,] -0.01321679 31.2201662
## [26,] 0.50335166 105.4845009
## [27,] 0.01246794 30.8852322
## [28,] 0.38044295 84.7188558
## [29,] 0.44284367 129.5256865
## [30,] 0.11488811 42.7695473

# Repeat the above simulation 1000 times, and store each matrix in a list
simulated_datasets <- map(1:1000, ~MASS::mvrnorm(30, mu=means, Sigma=covariance))

# Take the correlation of all the simulated matrices
simulated_correlation_matrices <- map(simulated_datasets, ~cor(.x))

# Take the average of all the simulated correlation matrices by summing them up
# and dividing by the total number of them
mean_correlation_matrix <- reduce(simulated_correlation_matrices, `+`) / length(simulated_correlation_matrices)

# View the mean correlation matrix from the 1000 simulations
mean_correlation_matrix
```

```
##          rec_att  rec_yds  rec_tds  rush_att  rush_yds  rush_tds  fumbles
## rec_att  1.0000000 0.9913964 0.9757739 0.4185241 0.4150325 0.2879686 0.8568449
## rec_yds  0.9913964 1.0000000 0.9869639 0.4033824 0.3994370 0.2817037 0.8550508
## rec_tds  0.9757739 0.9869639 1.0000000 0.3897496 0.3853983 0.2710301 0.8407170
## rush_att 0.4185241 0.4033824 0.3897496 1.0000000 0.9804176 0.8453839 0.4440133
## rush_yds 0.4150325 0.3994370 0.3853983 0.9804176 1.0000000 0.8585033 0.4379914
## rush_tds 0.2879686 0.2817037 0.2710301 0.8453839 0.8585033 1.0000000 0.2692011
## fumbles  0.8568449 0.8550508 0.8407170 0.4440133 0.4379914 0.2692011 1.0000000
## fpts      0.9895770 0.9981394 0.9916929 0.4385167 0.4354573 0.3186730 0.8540064
##          fpts
## rec_att  0.9895770
## rec_yds  0.9981394
## rec_tds  0.9916929
## rush_att 0.4385167
## rush_yds 0.4354573
## rush_tds 0.3186730
## fumbles  0.8540064
## fpts      1.0000000
```

### Question 3

21 points

Here's some code:

```
nDist <- function(n = 100) {
  df <- 10
  prob <- 1/3
  shape <- 1
  size <- 16
  list(
    beta = rbeta(n, shape1 = 5, shape2 = 45),
    binomial = rbinom(n, size, prob),
    chisquared = rchisq(n, df),
    exponential = rexp(n),
    f = rf(n, df1 = 11, df2 = 17),
```

```

    gamma = rgamma(n, shape),
    geometric = rgeom(n, prob),
    hypergeometric = rhyper(n, m = 50, n = 100, k = 8),
    lognormal = rlnorm(n),
    negbinomial = rnbinom(n, size, prob),
    normal = rnorm(n),
    poisson = rpois(n, lambda = 25),
    t = rt(n, df),
    uniform = runif(n),
    weibull = rweibull(n, shape)
  )
}

```

1. What does this do? (3 points)

```
round(sapply(nDist(500), mean), 2)
```

##	beta	binomial	chisquared	exponential	f
##	0.10	5.21	9.95	0.95	1.17
##	gamma	geometric	hypergeometric	lognormal	negbinomial
##	0.96	2.09	2.56	1.66	31.72
##	normal	poisson	t	uniform	weibull
##	0.02	25.20	0.04	0.46	1.06

The above code takes samples of size 500 from the distributions defined in `nDist()` and computes the mean

1. What about this? (3 points)

```
sort(apply(replicate(20, round(sapply(nDist(10000), mean), 2)), 1, sd))
```

##	beta	uniform	weibull	f	gamma
##	0.000000000	0.003077935	0.008944272	0.009233805	0.009514532
##	normal	exponential	t	hypergeometric	lognormal
##	0.010399899	0.010711528	0.011964861	0.014317821	0.019761739
##	binomial	geometric	poisson	chisquared	negbinomial
##	0.022542358	0.027028250	0.039483508	0.041877013	0.108171404

The above code takes samples of size 10000 from the distributions defined in `nDist()` and computes the mean

In the output above, a small value would indicate that `N=10,000` would provide a sufficient sample size

The below code estimates the sufficient sample size to estimate the mean for each distribution.

For each distribution, the procedure described above is computed starting from a sample size of 2, and the standard deviation of the 20 means is then stored. Until the standard deviation of the 20 means is below 0.02, the while loop continues, increasing the sample size each time. The first sample size to produce a standard deviation less than 0.02 is a rough estimate of the sufficient sample size to estimate the mean of the distribution.

```

df <- 10
prob <- 1/3
shape <- 1
size <- 16

l <- list(
  beta = function(n) rbeta(n, shape1 = 5, shape2 = 45),
  binomial = function(n) rbinom(n, size, prob),
  chisquared = function(n) rchisq(n, df),
  exponential = function(n) rexp(n),

```

```

f = function(n) rf(n, df1 = 11, df2 = 17),
gamma = function(n) rgamma(n, shape),
geometric = function(n) rgeom(n, prob),
hypergeometric = function(n) rhyper(n, m = 50, n = 100, k = 8),
lognormal = function(n) rlnorm(n),
negbinomial = function(n) rnbinom(n, size, prob),
normal = function(n) rnorm(n),
poisson = function(n) rpois(n, lambda = 25),
t = function(n) rt(n, df),
uniform = function(n) runif(n),
weibull = function(n) rweibull(n, shape)
)

```

```

#####
# Simulation to estimate the sufficient sample size
# to estimate the mean of each distribution
#
# This code may take a few minutes to run!
#####
set.seed(31415)

for (func in l){ # loop over all the distributions
  n <- 2
  stdev <- 1
  while (stdev > 0.02) # Loop until desired precision
  {
    # Simulate the data and take the standard deviation
    stdev <- map_dbl(1:20, ~func(n) %>% mean()) %>% sd()

    # Adaptively adjust the step size
    if(n < 500){n <- n+1}
    else if(n < 1000){n <- n+10}
    else if(n < 10000){n <- n+50}
    else{n <- n+100}
  }
  print(func)
  print(n)
  print(stdev)
}

```

```

## function(n) rbeta(n, shape1 = 5, shape2 = 45)
## <bytecode: 0x7fa3504f95b8>
## [1] 6
## [1] 0.01993172
## function(n) rbinom(n, size, prob)
## <bytecode: 0x7fa367dca0a8>
## [1] 4150
## [1] 0.0153046
## function(n) rchisq(n, df)
## <bytecode: 0x7fa355ad4d50>
## [1] 20200
## [1] 0.01658918
## function(n) rexp(n)
## <bytecode: 0x7fa367876978>

```

```

## [1] 2000
## [1] 0.01780965
## function(n) rf(n, df1 = 11, df2 = 17)
## <bytecode: 0x7fa347e6f9f0>
## [1] 403
## [1] 0.01931758
## function(n) rgamma(n, shape)
## <bytecode: 0x7fa361ded8b0>
## [1] 1700
## [1] 0.01936692
## function(n) rgeom(n, prob)
## <bytecode: 0x7fa367e13b30>
## [1] 8400
## [1] 0.01764024
## function(n) rhyper(n, m = 50, n = 100, k = 8)
## <bytecode: 0x7fa351e932b8>
## [1] 2000
## [1] 0.01976113
## function(n) rlnorm(n)
## <bytecode: 0x7fa3558b5328>
## [1] 8450
## [1] 0.01937498
## function(n) rnbinom(n, size, prob)
## <bytecode: 0x7fa355fc7358>
## [1] 73900
## [1] 0.01789058
## function(n) rnorm(n)
## <bytecode: 0x7fa3605614e8>
## [1] 1300
## [1] 0.01866401
## function(n) rpois(n, lambda = 25)
## <bytecode: 0x7fa35066a038>
## [1] 33600
## [1] 0.01880582
## function(n) rt(n, df)
## <bytecode: 0x7fa36045ffc8>
## [1] 2050
## [1] 0.01923908
## function(n) runif(n)
## <bytecode: 0x7fa353b71cb8>
## [1] 129
## [1] 0.01993969
## function(n) rweibull(n, shape)
## <bytecode: 0x7fa354d17b10>
## [1] 1600
## [1] 0.01971699

```

1. For each distribution, estimate the sample size required to simulate the distribution's mean. (15 points)

Don't worry about being exact. It should already be clear that  $N < 10,000$  for many of the distributions. You don't have to show your work. Put your answer to the right of the vertical bars (|) below.

distribution	N
beta	5
binomial	5,000

distribution	N
chisquared	20,000
exponential	2,000
f	500
gamma	2,000
geometric	8,000
hypergeometric	2,000
lognormal	10,000
negbinomial	80,000
normal	2,000
poisson	35,000
t	2,000
uniform	100
weibull	1,500